


# Monitoring Kafka Without Instrumentation Using eBPF

---

Ryan Cheng & Anton Rodriguez

# Our experience with Kafka

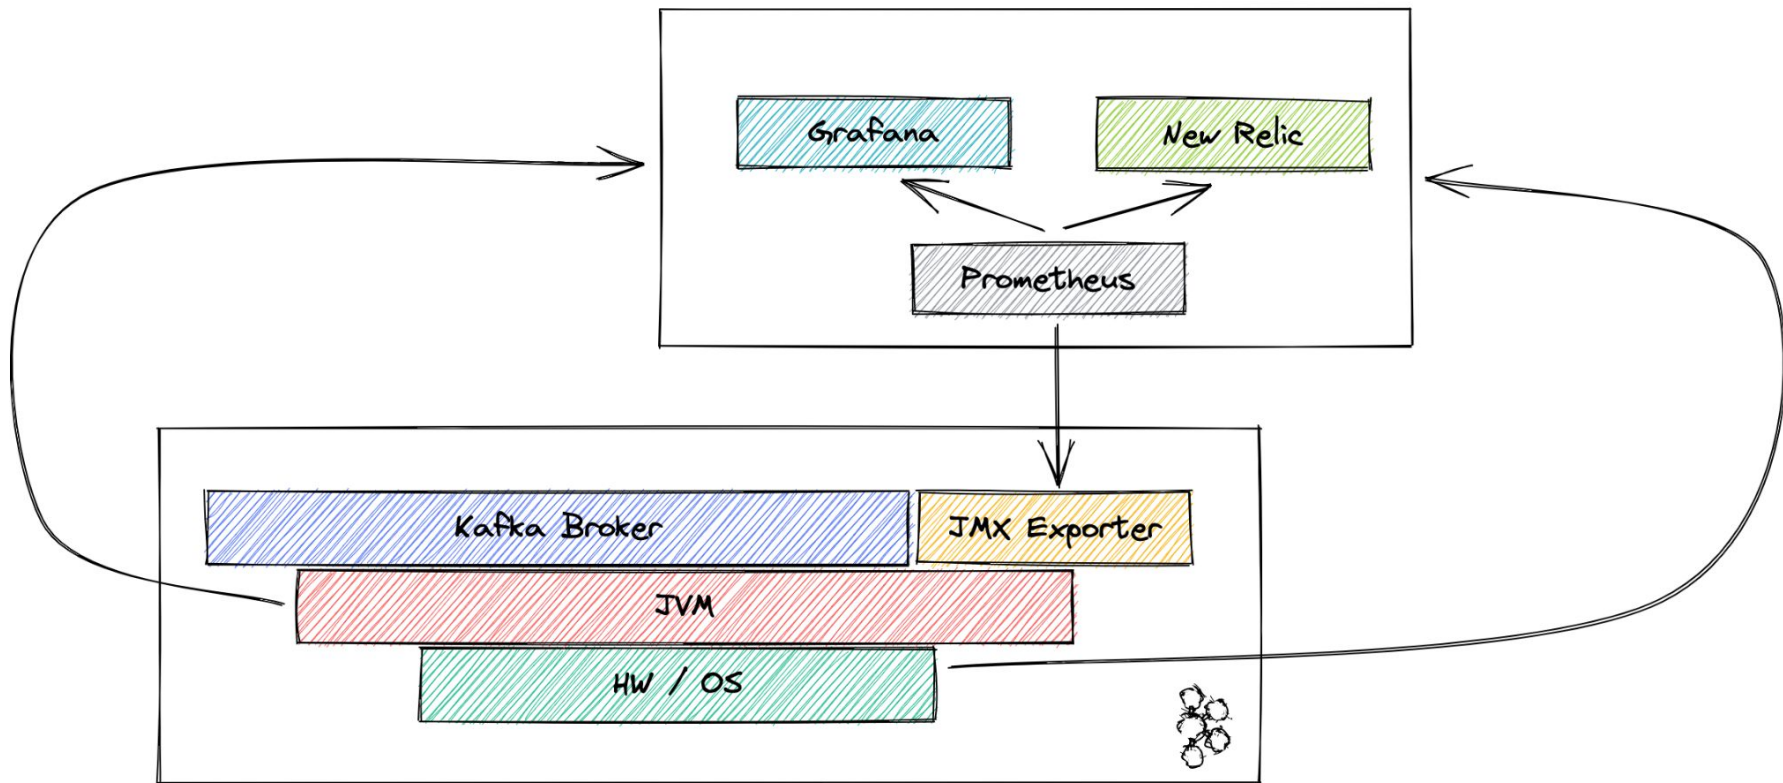


**125 PiB** of data per month


**3 B** data points per minute

production cluster of  
**~275 brokers** was  
running at **20 GB/sec**

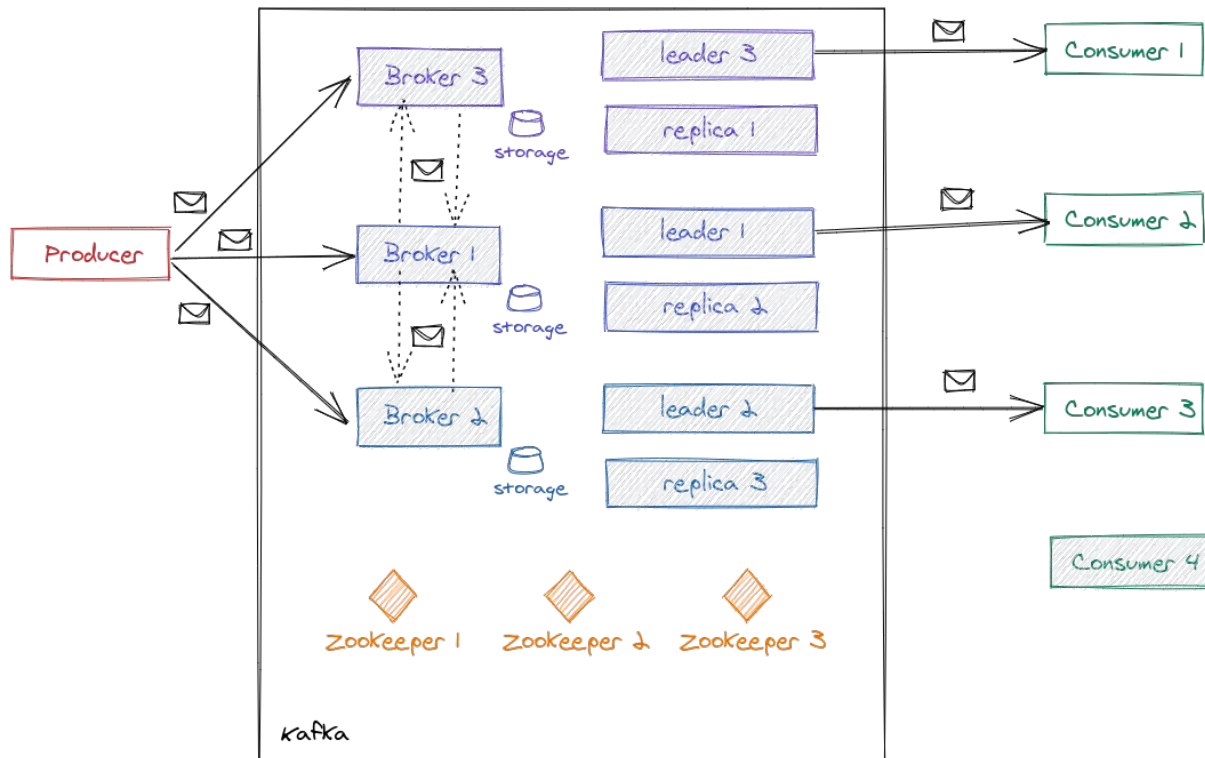
# Kafka & observability




# Kafka ecosystem



# Kafka rebalances

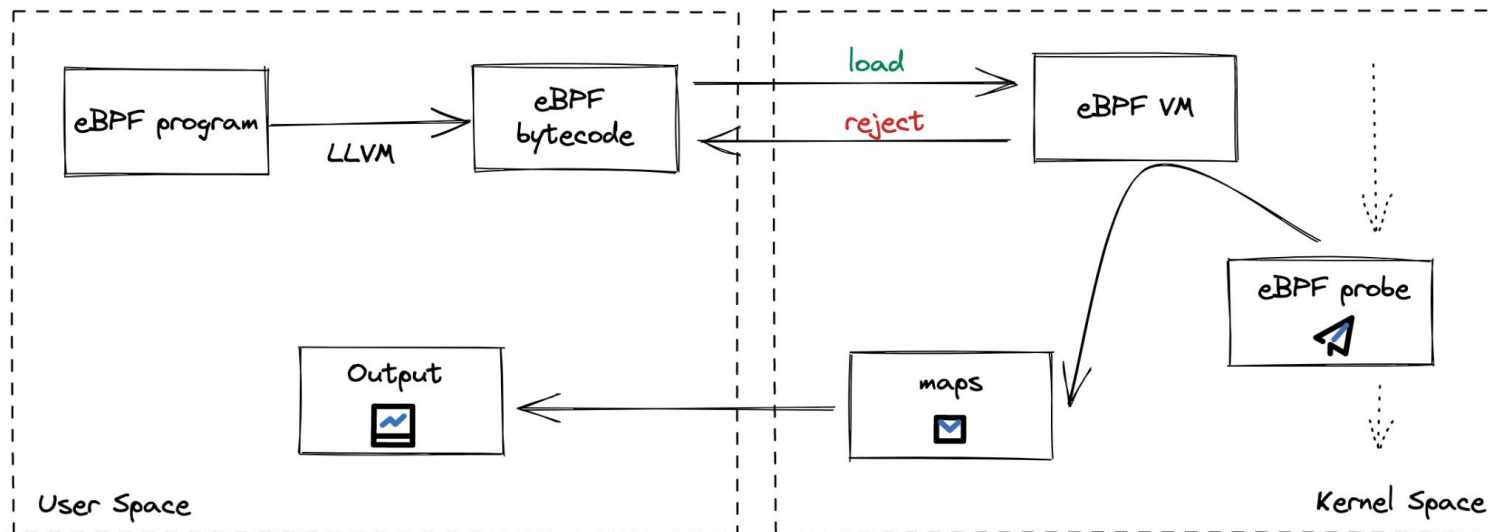


# Kafka Consumer Lag



# eBPF

*eBPF programs (user-defined, sandboxed bytecode executed by the kernel) allow user-defined instrumentation on a live kernel image that can never crash, hang or interfere with the kernel negatively*



# eBPF & BCC

## DISK I/O LATENCY HISTOGRAM

```
# biolatency
Tracing block device I/O... Hit Ctrl-C to end.
^C
```

usecs	: count	distribution
0 -> 1	: 0	
2 -> 3	: 0	
4 -> 7	: 0	
8 -> 15	: 0	
16 -> 31	: 0	
32 -> 63	: 0	
64 -> 127	: 1	
128 -> 255	: 12	*****
256 -> 511	: 15	*****
512 -> 1023	: 43	*****
1024 -> 2047	: 52	*****
2048 -> 4095	: 47	*****
4096 -> 8191	: 52	*****
8192 -> 16383	: 36	*****
16384 -> 32767	: 15	*****
32768 -> 65535	: 2	*

## TCP CONNECTIONS

```
# tcpconnect
PID  COMM      IP  SADDR          DADDR          DPORT
1479  telnet     4   127.0.0.1      127.0.0.1      23
1469  curl       4   10.201.219.236 54.245.105.25  80
1469  curl       4   10.201.219.236 54.67.101.145  80
1991  telnet     6   ::1             ::1             23
2015  ssh        6   fe80::2000:bff:fe82:3ac fe80::2000:bff:fe82:3ac 22
```

## NEW PROCESSES

```
# execsnoop
PCOMM  PID  RET  ARGS
bash   15887 0    /usr/bin/man ls
preconv 15894 0    /usr/bin/preconv -e UTF-8
man     15896 0    /usr/bin/tbl
man     15897 0    /usr/bin/nroff -mandoc -rLL=169n -rLT=169n -Tutf8
man     15898 0    /usr/bin/pager -s
nroff   15900 0    /usr/bin/locale charmap
nroff   15901 0    /usr/bin/groff -mtty-char -Tutf8 -mandoc -rLL=169n -rLT=169n
groff   15902 0    /usr/bin/troff -mtty-char -mandoc -rLL=169n -rLT=169n -Tutf8
groff   15903 0    /usr/bin/groff
```

## CUSTOM TRACING

```
# trace -p 2740 'do_sys_open "%s", arg2'
TIME  PID  COMM  FUNC  -
05:36:16 15872 ls    do_sys_open /etc/ld.so.cache
05:36:16 15872 ls    do_sys_open /lib64/libselinux.so.1
05:36:16 15872 ls    do_sys_open /lib64/libcap.so.2
05:36:16 15872 ls    do_sys_open /lib64/libacl.so.1
05:36:16 15872 ls    do_sys_open /lib64/libc.so.6
05:36:16 15872 ls    do_sys_open /lib64/libpcre.so.1
05:36:16 15872 ls    do_sys_open /lib64/libdl.so.2
05:36:16 15872 ls    do_sys_open /lib64/libattr.so.1
05:36:16 15872 ls    do_sys_open /lib64/libpthread.so.0
05:36:16 15872 ls    do_sys_open /usr/lib/locale/locale-archive
05:36:16 15872 ls    do_sys_open /home/vagrant
```



# bpfttrace

# kubectl-trace

## One-Liners

The following one-liners demonstrate different capabilities:

```
# Files opened by process
bpfttrace -e 'tracepoint:syscalls:sys_enter_open { printf("%s %s\n", comm, str(args->filename)); }'

# Syscall count by program
bpfttrace -e 'tracepoint:raw_syscalls:sys_enter { @[comm] = count(); }'

# Read bytes by process:
bpfttrace -e 'tracepoint:syscalls:sys_exit_read /args->ret/ { @[comm] = sum(args->ret); }'

# Read size distribution by process:
bpfttrace -e 'tracepoint:syscalls:sys_exit_read { @[comm] = hist(args->ret); }'

# Show per-second syscall rates:
bpfttrace -e 'tracepoint:raw_syscalls:sys_enter { @ = count(); } interval:s:1 { print(@); clear(@); }'

# Trace disk size by process
bpfttrace -e 'tracepoint:block:block_rq_issue { printf("%d %s %d\n", pid, comm, args->bytes); }'

# Count page faults by process
bpfttrace -e 'software:faults:1 { @[comm] = count(); }'

# Count LLC cache misses by process name and PID (uses PMCs):
bpfttrace -e 'hardware:cache-misses:1000000 { @[comm, pid] = count(); }'

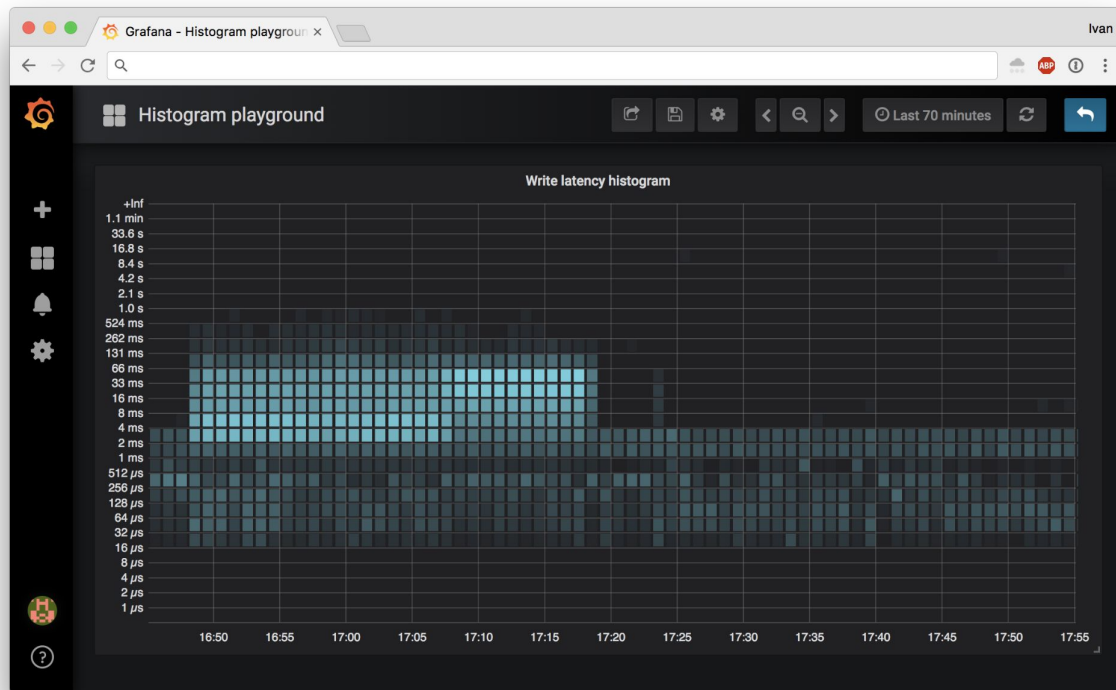
# Profile user-level stacks at 99 Hertz, for PID 189:
bpfttrace -e 'profile:hz:99 /pid == 189/ { @[ustack] = count(); }'

# Files opened, for processes in the root cgroup-v2
bpfttrace -e 'tracepoint:syscalls:sys_enter_openat /cgroup == cgroupid("/sys/fs/cgroup/unified/mycg")'
```

<https://github.com/iovisor/bpfttrace>

<https://github.com/iovisor/kubectl-trace>

# eBPF Exporter




[https://github.com/cloudflare/ebpf\\_exporter](https://github.com/cloudflare/ebpf_exporter)

# What's Pixie?


- **Open-source CNCF observability platform**
  - Using eBPF
- **Automatically traces network messages**
  - Kafka, HTTP, MySQL, etc.
  - Always active
- **No instrumentation**
  - No code modifications
  - No redeployments

A CNCF sandbox project



# Pixie's approach

- Pixie Edge Module (PEM) deploys on every node
- Capture data with eBPF
- Process data in user-space (protocol parsing)
- Store data into tables for querying by user

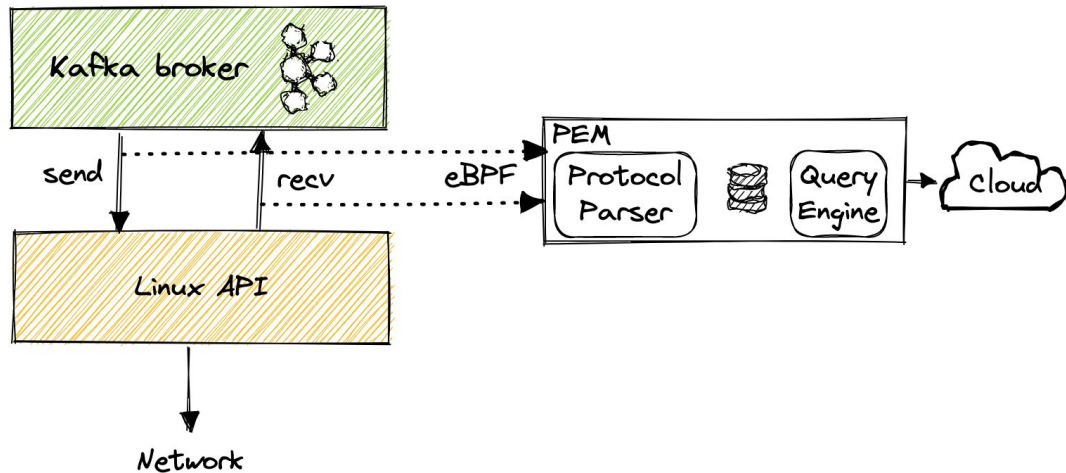


# How is the data traced?


Pixie traces network-related Linux syscalls with eBPF kprobes.

In the case of Kafka:

- Fetch and Produce messages
- JoinGroup and SyncGroup messages etc.



# Demo time!



```
bash -c "$(curl -fsSL https://withpixie.ai/install.sh)"
```

# Summary

- Kafka observability is challenging
- eBPF opens a new world of possibilities
- Pixie provides auto-instrumentation for Kubernetes applications
  - No code modification
  - No redeployment
  - Easy to use
  - Specific domain metrics
  - Low overhead

# Thank you!

## Questions?

