




Why is my eBPF code slow?

Simar Singh

Twitter: @simarpreet7

GitHub: @simar7

\$ whoami

1. I work on Open Source stuff at Aqua Security 
2. I like graphs & numbers 
3. I also like to grow plants 

We've all been here



We've all been here



- *"This code is slow"*

We've all been here



- *“This code is slow”*
- *“Need to improve performance”*

We've all been here



- *“This code is slow”*
- *“Need to improve performance”*
- *“Can we optimize this?”*

pop quiz: What is taking up CPU%?



Pick your best guess:

- Lot of computation?
- Waiting on something?

pop quiz: What is taking up CPU%?



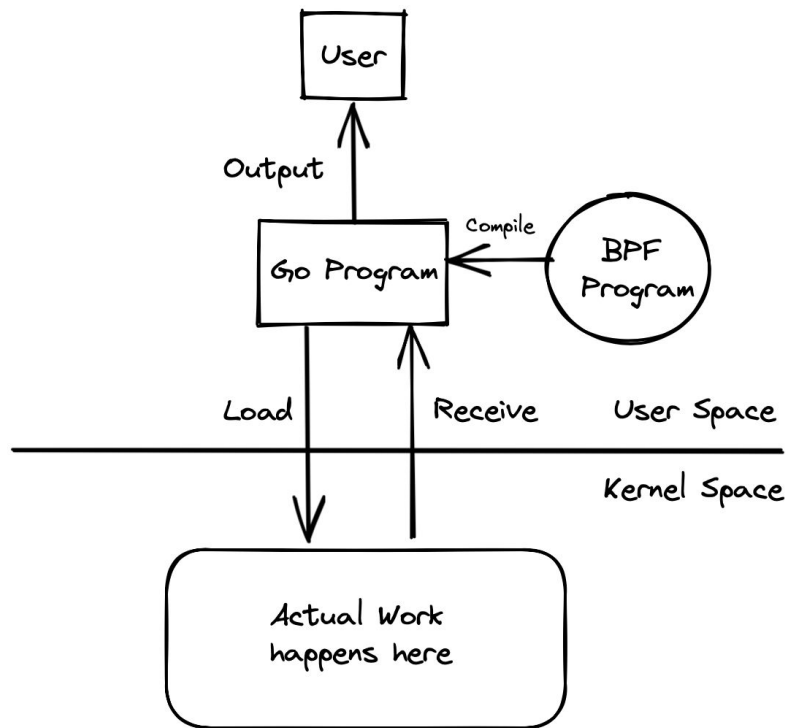
Pick your best guess:

- Lot of computation?
- Waiting on something?
- No idea.

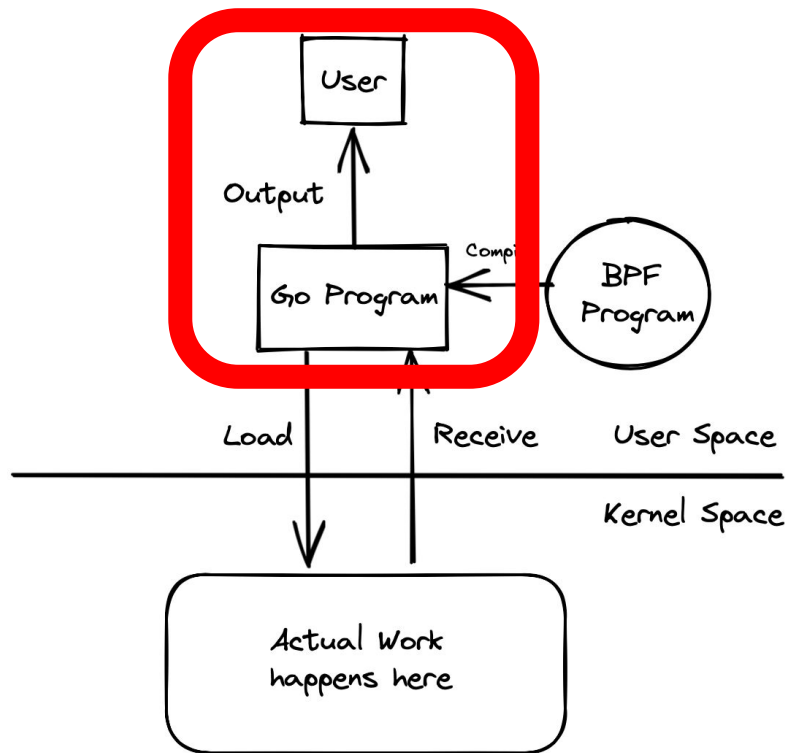
Usual benchmarking steps

- **Set a baseline**
 - What are the numbers today that matter?
- ***Kaizen* 改善: Continuous improvement**
 - Keep making small but incremental improvements
- **Know your limits**
 - No software is perfect, it's always a tradeoff


The stack




The stack



How to set a baseline




How to set a baseline




Measuring from the outside looking in

How to set a baseline





<https://github.com/simar7/simple-linux-monitoring-stack>



Measuring from the outside looking in

How to set a baseline


<https://github.com/simar7/simple-linux-monitoring-stack>



Pros	Cons
10,000ft overview	Hard to pinpoint
Easy to setup	Many components at play


Measuring from the outside looking in

How to set a baseline



Measuring from within the system

How to set a baseline



<https://github.com/pkg/profile>

```
func main() {  
+   defer profile.Start(profile.CPUProfile, profile.NoShutdownHook, profile.ProfilePath(".")).Stop()  
+   app := &cli.App{
```

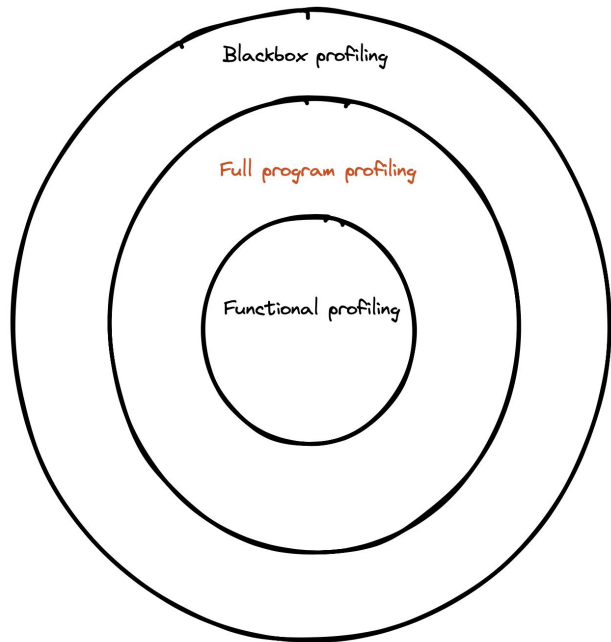
```
→ ~ go tool pprof -http=:6060 cpu.pprof  
Serving web UI on http://localhost:6060
```

1.48s

runtime
cgocall
1.48s (77.08%)

Measuring from within the system

How to set a baseline



Measuring from within the system

<https://github.com/pkg/profile>


```
func main() {  
+   defer profile.Start(profile.CPUProfile, profile.NoShutdownHook, profile.ProfilePath(".")).Stop()  
+  
+   app := <cli.App>  
}
```

```
→ ~ go tool pprof -http=:6060 cpu.pprof  
Serving web UI on http://localhost:6060
```

1.48s
↓
runtime
cgocall
1.48s (77.08%)


Pros	Cons
Easy to setup	No visibility outside of userspace
Code paths highlighted	Can have low Signal to Noise ratio

How to set a baseline



Measuring each unit on it's own


How to set a baseline



Measuring each unit on it's own

```
func BenchmarkFoo(b *testing.B){  
    setup()  
    b.ResetTimer()  
    for i := 0; i < b.N; i++ {  
        doWork()  
    }  
}
```

How to set a baseline

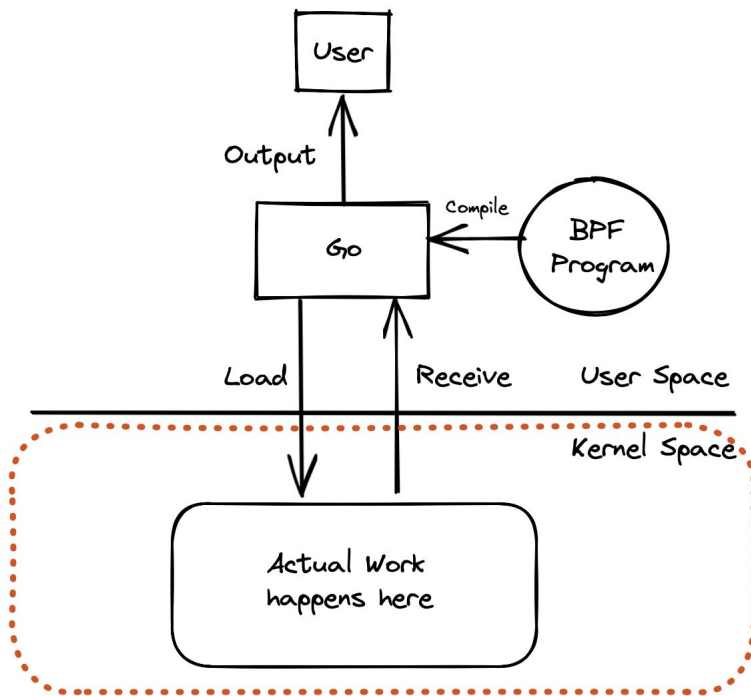


Measuring each unit on it's own

```
func BenchmarkFoo(b *testing.B){
    setup()
    b.ResetTimer()
    for i := 0; i < b.N; i++ {
        doWork()
    }
}
```

Pros	Cons
Small unit of work	Can lack external attributing factors
Well defined	Improvements might not make a big impact

What about the kernel space code?



eBPF to the rescue!

```
// Taken from iovisor/bcc/tools/offcputime.py

// record previous thread sleep time
if ((THREAD_FILTER) && (STATE_FILTER)) {
    ts = bpf_ktime_get_ns();
    start.update(&pid, &ts);
}
// get the current thread's start time
pid = bpf_get_current_pid_tgid();
tgid = bpf_get_current_pid_tgid() >> 32;
tsp = start.lookup(&pid);
if (tsp == 0) {
    return 0;          // missed start or filtered
}

// calculate delta
...
```

⚡ eBPF to the rescue!

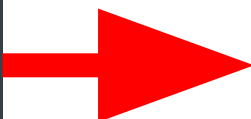
```
// Taken from iovisor/bcc/tools/offcputime.py

// record previous thread sleep time
if ((THREAD_FILTER) && (STATE_FILTER)) {
    ts = bpf_ktime_get_ns();
    start.update(&pid, &ts);
}

// get the current thread's start time
pid = bpf_get_current_pid_tgid();
tgid = bpf_get_current_pid_tgid() >> 32;
tsp = start.lookup(&pid);
if (tsp == 0) {
    return 0;          // missed start or filtered
}

// calculate delta
...
```

<https://github.com/iovisor/bcc/blob/master/tools/offcputime.py>



```
finish_task_switch
schedule
schedule_hrttimeout_range_clock
schedule_hrttimeout_range
ep_poll
do_epoll_wait
__x64_sys_epoll_wait
do_syscall_64
entry_SYSCALL_64_after_hwframe
-          myprogram (119161)
        63564745

finish_task_switch
schedule
futex_wait_queue_me
futex_wait
do_futex
__x64_sys_futex
do_syscall_64
entry_SYSCALL_64_after_hwframe
-          myprogram (119164)
        65074899

finish_task_switch
schedule
schedule_hrttimeout_range_clock
schedule_hrttimeout_range
ep_poll
do_epoll_wait
__x64_sys_epoll_wait
do_syscall_64
entry_SYSCALL_64_after_hwframe
-          myprogram (119165)
        69912863
```


⚡ eBPF to the rescue!


```
// Taken from iovisor/bcc/tools/offcputime.py

// record previous thread sleep time
if ((THREAD_FILTER) && (STATE_FILTER)) {
    ts = bpf_ktime_get_ns();
    start.update(&pid, &ts);
}

// get the current thread's start time
pid = bpf_get_current_pid_tgid();
tgid = bpf_get_current_pid_tgid() >> 32;
tsp = start.lookup(&pid);
if (tsp == 0) {
    return 0;          // missed start or filtered
}

// calculate delta
...
```

<https://github.com/iovisor/bcc/blob/master/tools/offcputime.py>



```
finish_task_switch
schedule
schedule_hrttimeout_range_clock
schedule_hrttimeout_range
ep_poll
do_epoll_wait
__x64_sys_epoll_wait
do_syscall_64
entry_SYSCALL_64_after_hwframe
- 63564745 myprogram (119161)

finish_task_switch
schedule
futex_wait_queue_me
futex_wait
do_futex
__x64_sys_futex
do_syscall_64
entry_SYSCALL_64_after_hwframe
- 65074899 myprogram (119164)

finish_task_switch
schedule
schedule_hrttimeout_range_clock
schedule_hrttimeout_range
ep_poll
do_epoll_wait
__x64_sys_epoll_wait
do_syscall_64
entry_SYSCALL_64_after_hwframe
- 69912863 myprogram (119165)
```

pop quiz: revisited



Pick your best guess:

- Lot of computation?
- Waiting on something?


pop quiz: revisited



Pick your best guess:

- Lot of computation?
- Waiting on something?

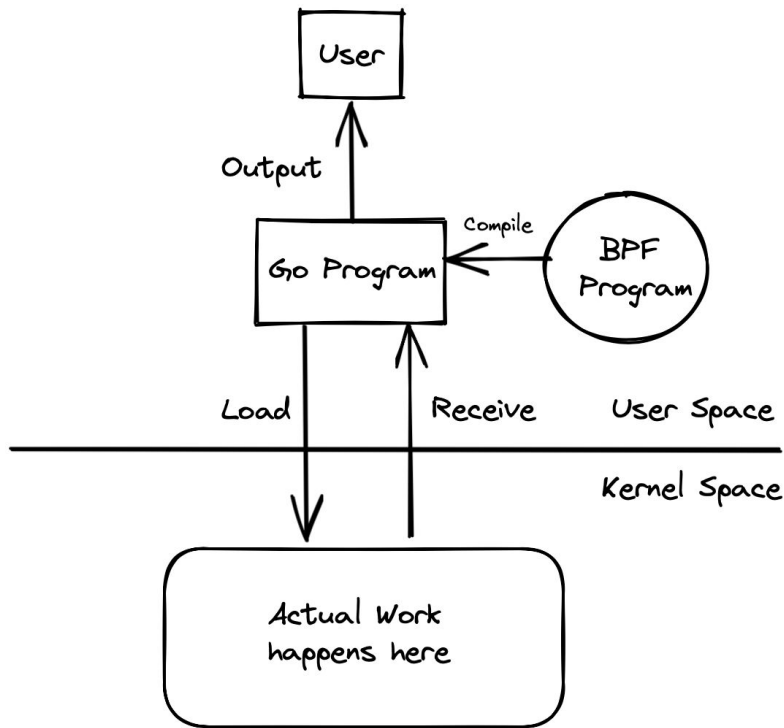
pop quiz: revisited



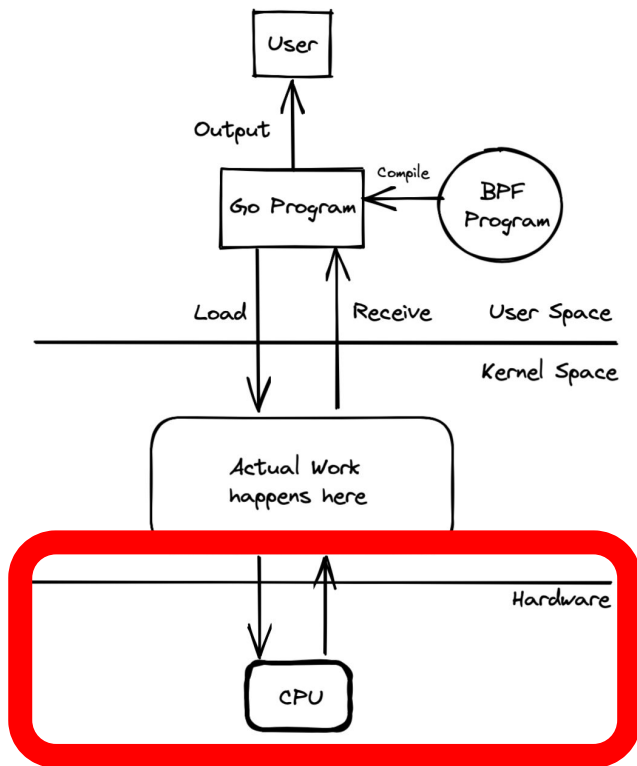
Pick your best guess:

- Lot of computation?
- **Waiting on something?**

Can we go even lower?



Can we go even lower? yes.



Visiting an old friend

perf (Linux)

From Wikipedia, the free encyclopedia

perf (sometimes called **perf_events**^[1] or **perf tools**, originally **Performance Counters for Linux, PCL**)^[2] is a [performance analyzing](#) tool in [Linux](#), available from [Linux kernel](#) version 2.6.31 in 2009.^[3] [Userspace](#) controlling utility, named `perf`, is accessed from the [command line](#) and provides a number of [subcommands](#); it is capable of statistical profiling of the entire system (both kernel and userland code).

It supports [hardware performance counters](#), [tracepoints](#), software performance counters (e.g. hrtimer), and dynamic engineers recognized perf (along with [OProfile](#)) as one of the Linux.^[5]

Visiting an old friend

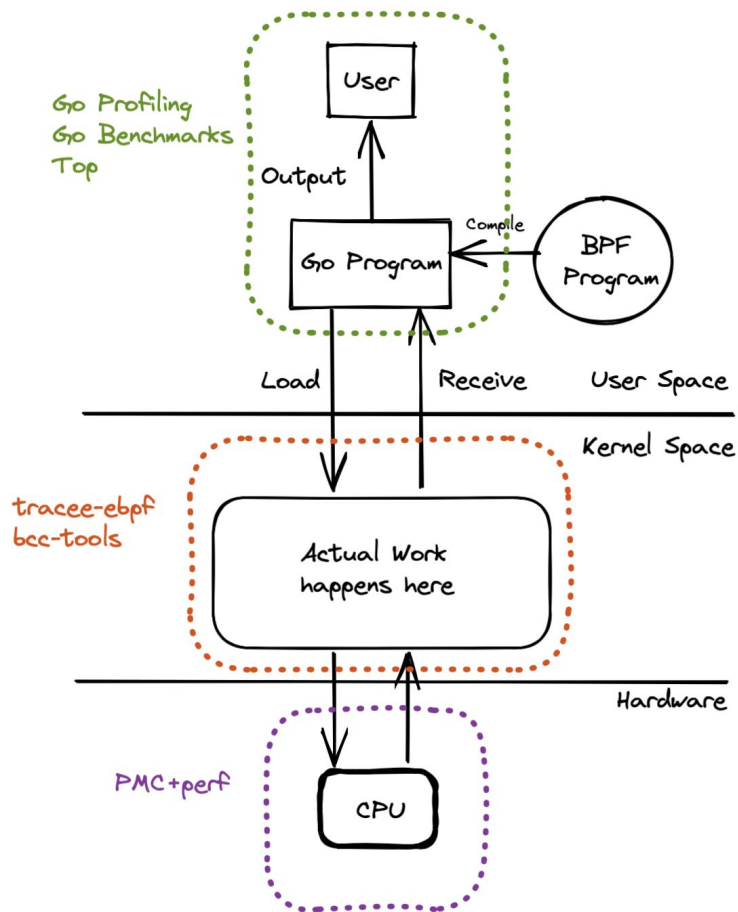
```
$ perf stat -a -- ./myprogram
```

```
Performance counter stats for 'system wide':
```

139,679.29 msec	cpu-clock	#	4.000 CPUs utilized
27,902	context-switches	#	0.200 K/sec
231	cpu-migrations	#	0.002 K/sec
7,616	page-faults	#	0.055 K/sec
11,254,036,556	cycles	#	0.081 GHz
5,513,412,312	instructions	#	0.48 insn per cycle
1,989,463,444	branches	#	14.243 M/sec
22,198,708	branch-misses	#	1.12% of all branches

```
34.921089105 seconds time elapsed
```


Wrap up



Thanks!

Simar Singh

Twitter: @simarpreet7

GitHub: @simar7