

Zach Wasserman – CTO, Fleet




eBPF and the future of osquery on Linux

osquery@scale 2021



Who am I

- Co-creator of osquery
- Cofounder & CTO of FleetDM
- (Former) Cofounder & Principal Engineer of Kolide




Silent Accomplice

Alessandro Gario

Senior Software Engineer

Trail of Bits



Audit & Osquery

Audit & Osquery

Tables

- Audit powers many of the event-based tables for osquery on Linux
 - process_events
 - process_file_events
 - socket_events
 - user_events
 - apparmor_events
 - selinux_events



Audit & Osquery

Configuration

- Base configuration to enable audit
 - `--disable_audit=false`
 - `--audit_allow_config=true`
 - `--audit_persist=true`



Audit & Osquery

Configuration

- Enable each feature separately
 - `--audit_allow_apparmor_events`
 - `--audit_allow_fim_events`
 - `--audit_allow_fork_process_events`
 - `--audit_allow_kill_process_events`
 - `--audit_allow_process_events`
 - `--audit_allow_selinux_events`
 - `--audit_allow_sockets`
 - `--audit_allow_user_events`



Audit Drawbacks



There can be only one

Audit Drawbacks

There can be only one

- Audit's design allows only a single consumer of the generated events
- Receiving audit events in osquery means disabling auditd
- Disable auditd -> No audit events written to file
 - Some tools expect to be able to retrieve audit events from file!



Audit Drawbacks

There can be only one

- Many SELinux tools rely on audit logs in `/var/log/audit/audit.log`
 - `sealert`
 - `audit2allow`



Audit Drawbacks

Containers

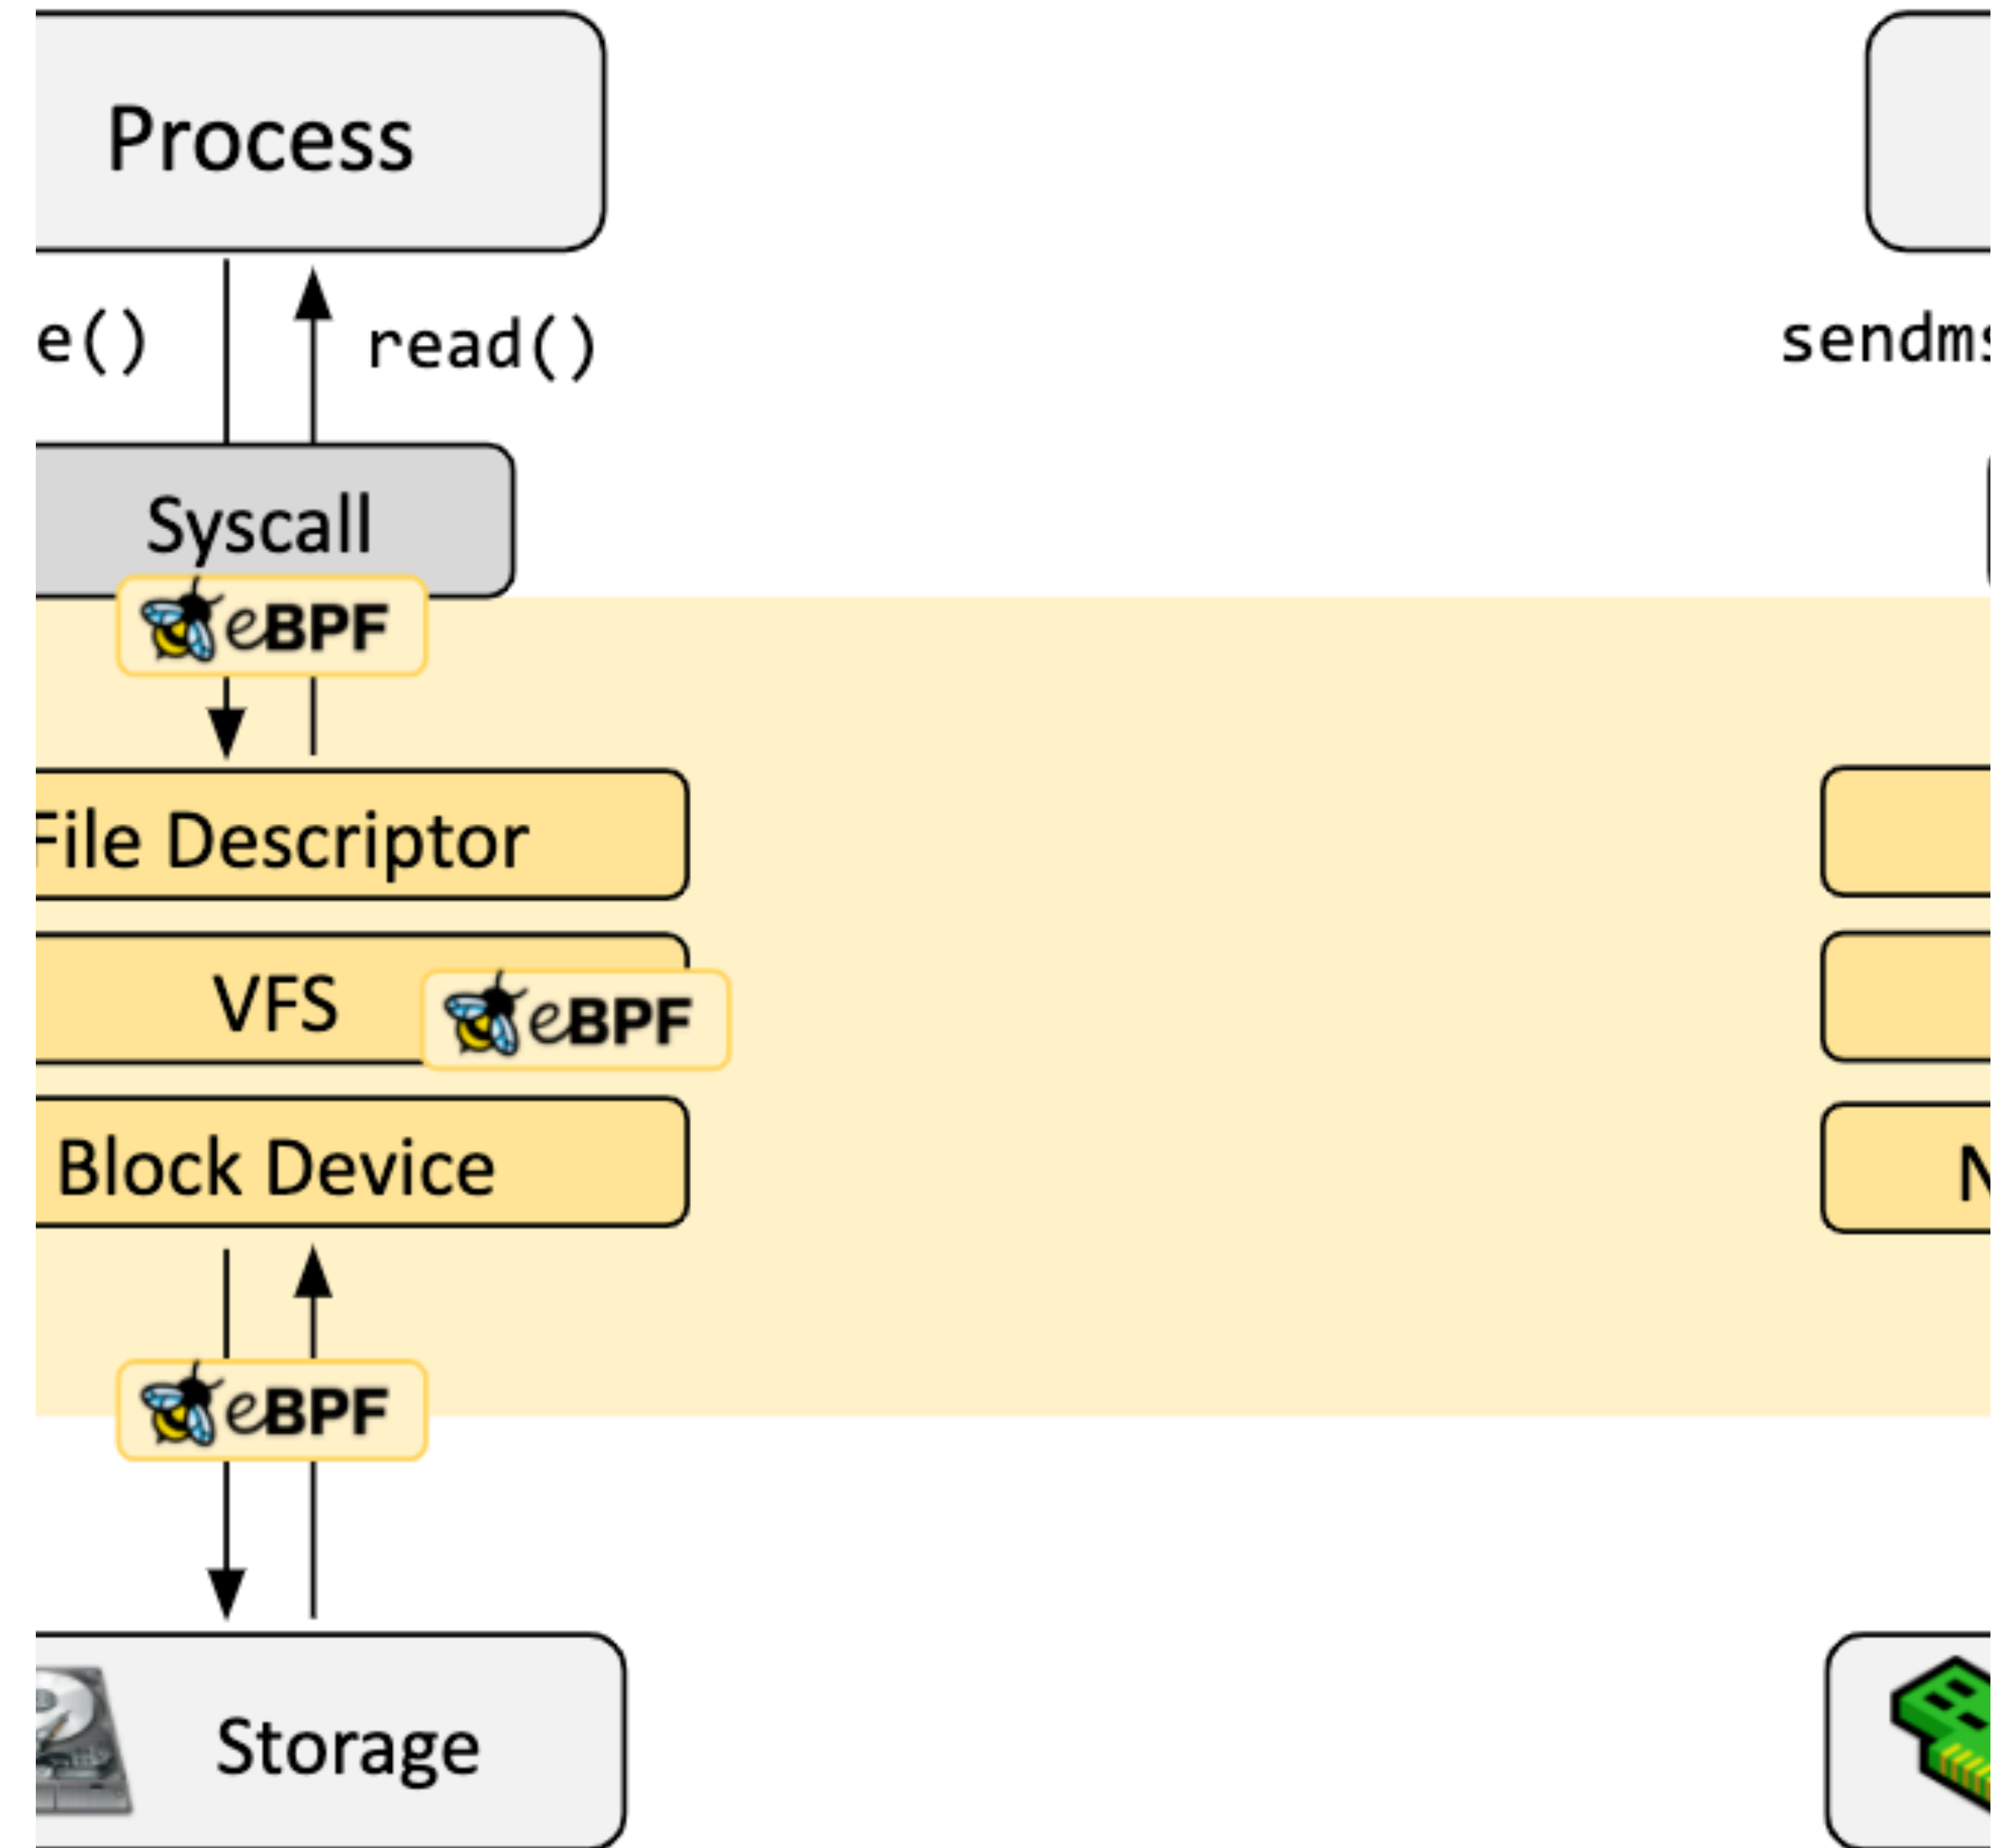
- Audit does actually work with containers
 - Fixed in Red Hat Bug 893751
- Audit is not “aware” of containers
 - Lack of namespace information hinders usability in container workloads



eBPF

eBPF

Programmable Hooks



eBPF

Safety

- Required privileges
- Program verification
 - Limited size
 - Limited complexity
 - Bounded loops
- Controlled memory access



eBPF & Osquery

eBPF & Osquery

State

- eBPF functionality released in osquery 4.6.0
 - Built primarily by Alessandro Gario
- Implemented on top of github.com/trailofbits/ebpfpub



eBPF & Osquery

Tables

- `bpf_process_events`
- `bpf_socket_events`
- ...



eBPF & Osquery

Configuration

- `--enable_bpf_events`
- That's it!



eBPF & Osquery

Tuning

- `--bpf_buffer_storage_size` (default 512)
- `--bpf_perf_event_array_exp` (default 10)



eBPF & Osquery

Compatibility

- Targeting Kernels 4.18+ (2018)
 - Possible to extend compatibility back to 4.10+ (2017)
- eBPF Probes are generated at runtime
 - One binary can work on most Kernels



eBPF & Osquery

Coming soon...

- Support for correlating BPF events with containers
 - Mapping cgroup_ids to Docker containers
- `process_dns_events`



eBPF & Osquery

Future

- We now have a pattern for instrumenting nearly anything on Linux
 - System calls
 - Kernel tracepoints
 - User-space function calls
- These can be dynamically configured at osquery runtime



eBPF & Osquery

Future - Security

- Instrument any and all syscalls of interest
- Track signals sent to processes
- Kernel module loads/unloads
- Track LD_PRELOAD



eBPF & Osquery

Future - Devops/SRE

- Instrument any and all syscalls of interest(!)
- Measure latency and resource consumption of OS processes
 - Network stack
 - Filesystem
 - Other I/O
- Count and measure functions within user-space



eBPF & Osquery

Future - Imagine

- Let's look at the tools of today
 - bpftrace
 - BCC
- Which of these use cases map well to osquery's SQL model?
- How can osquery be useful for shipping the aggregated information from hosts?



Conclusion

Audit & eBPF are both viable approaches

**eBPF has potential to dramatically increase
scope of observability with osquery**

Thank you
zach@fleetdm.com
🐦 @thezachw
⌘ @zwass

